

Robust & Asymptotically Locally Optimal UAV-Trajectory Generation Based on Spline Subdivision

Ruiqi Ni¹, Teseo Schneider², Daniele Panozzo², Zherong Pan³, Xifeng Gao¹

Abstract—Generating locally optimal UAV-trajectories is challenging due to the non-convex constraints of collision avoidance and actuation limits. We present the first local, optimization-based UAV-trajectory generator that simultaneously guarantees validity and asymptotic optimality for known environments. **Validity:** Given a feasible initial guess, our algorithm guarantees the satisfaction of all constraints throughout the process of optimization. **Asymptotic Optimality:** We use an asymptotic exact piecewise approximation of the trajectory with an automatically adjustable resolution of its discretization. The trajectory converges under refinement to the first-order stationary point of the exact non-convex programming problem. Our method has additional practical advantages including joint optimality in terms of trajectory and time-allocation, and robustness to challenging environments as demonstrated in our experiments.

I. INTRODUCTION

Unmanned Aerial Vehicle (UAV) finds many real-world applications in inspection, search and rescue, and logistic automation. A key challenge in safe UAV-trajectory planning is to find locally optimal flying strategies, in terms of energy/time efficiency and smoothness, while accounting for various dynamic and kinematic constraints. A valid trajectory needs to satisfy two constraints to be physically realizable: the actuation limits must be respected (dynamic constraint), and the robot should maintain a safety distance from the boundary of the freespace (kinematic constraint). These two constraints combined define a non-convex and non-smooth feasible domain, which is notoriously difficult to handle for optimization-based motion planners [39], [24], [28].

Prior works tackle the non-convex, non-smooth constraints in one of three ways: infeasible recovery, relaxation, and global search. **Infeasible Recovery:** Nonlinear constrained optimizers, e.g. Sequential Quadratic Programming (SQP), have been used in [39], [34], [30] to directly handle non-convex constraints. These methods rely on the constraints' directional derivatives and penalty functions to pull infeasible solutions back to the feasible domain. However, their feasibility is not guaranteed: an example of a typical scenario of infeasible solutions involving an environment with multiple, thin-shell obstacles is discussed in [28]. **Relaxation:** In [6], [3], relaxation schemes are used to limit the solutions to a (piecewise) convex and smooth subset. For example, the freespace is approximated by the union of convex subsets and a UAV-trajectory is represented by piecewise splines, where

This work was supported in part by NSF-IIS-(1910486 and 1908767).

¹Ruiqi Ni and Xifeng Gao are with Department of Computer Science, Florida State University (rn19g@my.fsu.edu and gao@cs.fsu.edu).

²Teseo Schneider and Daniele Panozzo are with the Department of Computer Science, New York University (teseo.schneider@nyu.edu and panozzo@nyu.edu). ³Zherong Pan is with the Department of Computer Science, University of Illinois Urbana-Champaign (zherong@illinois.edu).

each piece is constrained in one convex subset. These relaxed formulations can be efficiently handled by modern (mixed-integer) convex optimization tools. However, these methods can only find sub-optimal solutions because the convex subset is a strict inner-approximation with a non-vanishing gap (and increasing the approximation power increases the algorithmic running time). For environments with narrow passages, relaxation can even turn a feasible problem into an infeasible one. **Global Search:** Kinodynamic-RRT* [35] and its variants can find globally optimal UAV trajectories in a discrete space and stochastic gradient descend [13] will approach locally optimal solutions only if the descent direction is approximately by infinitely many samples. However, these algorithms take excessively large number of samples. To mitigate their computational cost, prior methods such as [19] have to terminate sampling early and then numerically rectify the solutions, and these numerical rectification methods still suffer from the limitations of infeasible recover or relaxation approaches.

Main Results: We propose a new formulation of local, optimization-based motion planning for UAV-trajectories with robustness and asymptotic optimality guarantee. Our method represents the UAV-trajectory using smooth curves, encoded as piecewise Bézier curves. For each Bézier piece, we formulate the collision-free constraints between its control polygon and the environmental geometry. We show that these constraints can be formulated as a summation of distance functions between geometric primitives (e.g. point-triangle and edge-edge pairs). Each distance function is differentiable when restricted to the feasible domain and the restriction can be achieved using a primal interior point method (P-IPM).

As compared with infeasible recovery methods that maintain both primal and dual variables, our primal-only formulation provides guaranteed solution feasibility throughout the entire process of optimization. This feature allows our planner to answer anytime re-planning queries as in [17] because the planner can be terminated at any iteration and return a feasible solution. And compared with relaxation-based methods, our method can asymptotically minimize the sub-optimality gap due to the control-polygon approximation of the true trajectory. This is achieved by using an adaptive Bézier subdivision scheme with user-controllable termination criterion.

Our method contributes to the solving of UAV planning problem in three ways: 1) Combined with RRT-style approaches, we can robustly optimize a feasible initial guess. 2) Starting from a feasible initial guess, our primal solver is guaranteed to preserve the homotopy class of the trajectory. This feature is useful, for example, for inspection planners

where the order of inspection cannot be altered. 3) Our formulation achieves simultaneous optimality in trajectory shapes and time allocations.

II. RELATED WORK

Trajectory optimization for a general robotic system can be solved using dynamic programming [29], direct collocation [33], and path-integral control [13], [37]. While widely applicable, these methods require a search in a high-dimensional space and are thus computationally demanding and sensible to numerical failures. For the special case of UAV dynamics, a more robust and efficient approach has been proposed in [20], where the algorithm optimizes a reference position trajectory (global search) and then recovers the control inputs making use of differential flatness (local optimization). Our method follows the same strategy but offers the additional advantage of feasibility guarantees and locally optimality of the original non-convex optimization problem.

Global Search provides an initial trajectory that routes the UAV from a start to a goal position and usually optimizes the trajectory by minimizing a state- or time-dependent objective function. Global search can be accomplished using sampling-based motion planner [36], [16], mixed-integer programming [3], discrete search [18], and fast-marching method [9]. All these methods are approximating the globally optimal UAV-trajectory under some assumptions: optimal sampling-based motion planner approaches the optimal solution after a sufficient number of samples have been drawn; mixed-integer programming and discrete search restricts the decision space to a disjoint-convex or discrete subset, respectively; fast-marching methods discretize both the trajectory and the environment on a uniform grid. Global search is essential not only for providing an initial guess to the local optimization, but also for satisfying additional constraints. Typical constraints include visitation order [10], homotopy class [14], and coverage [11], [5]. These constraints can be formulated into the mixed-integer programming via additional binary decision variables [3] or into discrete search algorithms such as A* by pruning trajectories that violate the constraints [18]. All these methods are orthogonal to our work and any global search method with the collision-free guarantee can be used as the initialization step of our method.

Local Optimization complements the global search by refining the trajectories in a neighborhood of the initial guess. This step lifts the restricted search space assumption in the global search and further minimizes the objective function in the continuous decision space, until a local optimal solution is attained. In addition, the local step ensures that the solution satisfies kinematics and dynamics constraints so that the trajectory is executable on hardware. All prior works formulate the local step as a continuous or discrete optimization problem. The continuous problem can only be solved in the obstacle-free cases with a closed-form solution [35], [18]. In [9], [31], [34], obstacle-free constraints are relaxed to be convex and the resulting discrete optimization is guaranteed to be solvable using primal-dual interior point method (PD-IPM). In [38], [8], obstacle-free constraints are

considered in its original, non-convex form, which is solved by PD-IPM but without feasibility guarantee. In this work, we show that the feasibility guarantee can be provided using P-IPM.

Time-Allocation is an essential part of local optimization formulated as part of the dynamic constraints. Early works [27], [9] first optimize the shape of the trajectory and then allocate time for each discrete segment, and the resulting trajectory is sub-optimal with respect to the time variables. It has been shown that, for a fixed trajectory, optimal time-allocation can be achieved using bang-bang solutions computed via numerical integration [25]. The trajectories computed using numerical integration by [25] are locally optimal in either shape variables or time variables but not both at the same time. The latest works [30], [31] achieve shape-time joint optimality via bilevel optimization or weighted combination. We follow [31] and optimize a weighted combination of shape and time cost functions.

III. PROBLEM STATEMENT

We propose a new method for local optimization of UAV trajectories. A UAV moves along the center-of-mass trajectory $p(t, W)$ with $t \in [0, T]$ where T is the travel time and W is the set of decision variables. A local optimizer takes an initial trajectory as input and locally updates it to minimize a cost function:

$$\mathcal{O}(W, T) = \int_0^T c(p(t, W), t) dt + R(p(t, W), T), \quad (1)$$

where $c(p, t)$ is the state-dependent cost (e.g., dependent on velocity, acceleration, jerk, or snap) and $R(p, T)$ is the terminal cost. In the meantime, a feasible UAV trajectory should satisfy a set of kinematic, dynamics, and application-dependent constraints. In this paper, we assume that dynamics constraints take the form of velocity and acceleration limits:

$$\|\dot{p}(t, W)\| \leq v_{max} \quad \|\ddot{p}(t, W)\| \leq a_{max} \quad \forall t \in [0, T], \quad (2)$$

which are semi-infinite constraints in the variable t . Here $\{v, a\}_{max}$ are the maximum allowable velocity and acceleration. We further assume only the collision-free kinematic constraint that can take different forms depending on how obstacles are represented. Prominent environmental representations are point clouds and triangle meshes. These two representations can be uniformly denoted as a tuple of discrete elements $\mathcal{E} = \langle \mathcal{P}, \mathcal{L}, \mathcal{T} \rangle$ where \mathcal{P} is a set of points, \mathcal{L} is a set of line segments, and \mathcal{T} is a set of triangles. The distance between a point x and the environment is denoted as $\text{dist}(x, \mathcal{E})$. Throughout the paper, we slightly abuse notation and define $\text{dist}(\bullet, \bullet)$ as closest distance between a pair of geometric objects (e.g. points, line segments, triangles, convex hulls, point clouds, and triangle meshes). Our complete local trajectory optimization problem can be written as:

$$\underset{W, T}{\operatorname{argmin}} \mathcal{O}(W, T) \quad \text{s.t. } \forall t \in [0, T] \begin{cases} \|\dot{p}(t, W)\| \leq v_{max} \\ \|\ddot{p}(t, W)\| \leq a_{max} \\ \text{dist}(p(t, W), \mathcal{E}) \geq d_0 \end{cases}, \quad (3)$$

where d_0 is an uncertainty-tolerating safety distance. Equation 3 is among the most challenging problem instances in

operations research due to the semi-infinite constraints in time t and the non-smoothness of the function $\text{dist}(\bullet, \bullet)$ in general [9], [12].

A. Trajectory Representation

We represent $p(t, W)$ as a composite Bézier curve where each Bézier piece has a constant degree M . We note that this generic definition allow for curves of arbitrary continuity. The continuity across the different Bézier pieces is achieved by constraining the control points. The constraints are linear and can thus be removed from the system based on Stärk's construction [26].

IV. SUBDIVISION-BASED P-IPM

We propose an iterative algorithm, inspired by the recently proposed incremental potential contact handling scheme [15], to find locally optimal solutions of Equation 3 with guaranteed feasibility.

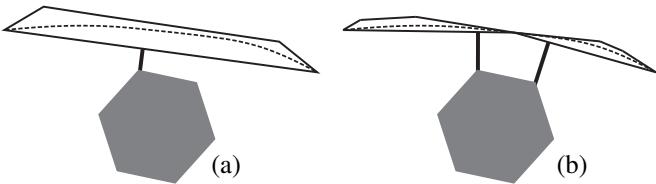


Fig. 1: (a): Our method imposes distance constraints between the convex hull of control points and the obstacle. (b): Our constraints can be refined by subdivision.

A. Method Overview

The key idea of [15] is to use a primal-only (instead of primal-dual) interior point method to solve collision-constrained nonlinear optimization problems. This is achieved by transforming the collision-constraints into log-barrier functions, and reducing the problem to non-linear, unconstrained optimization. Although primal-dual approaches have better numerical stability in many problems, their convergence relies on sufficient smoothness and Mangasarian-Fromovitz constraint qualification (MFCQ) along the central path, which do not hold for collision constraints, often leading to failures in feasibility recovery. In contrast, given a feasible initial guess, line-search based, primal-only solvers are guaranteed (even in floating point implementations), to stay inside the feasible domain. The key to establish this invariance is the use of a robust line-search scheme that prevents solution from leaving the feasible domain, in our case rejecting iterations with a jump in homotopy class or violation of collision-free constraints. It has been shown in [15] that, for linearized trajectories, these safety checks can be performed using linear continuous collision detection (CCD). In this work, we extend this construction for motion planning with curved trajectories.

A safe line-search scheme ensures feasibility but does not guarantee local optimality, i.e. convergence to a first-order critical point. Assuming gradient-based optimizers are used, the additional requirement for local optimality is the objective function Equation 3 being differentiable,

which is challenging due to the non-smoothness of the function $\text{dist}(\bullet, \bullet)$. Although SQP algorithms with gradient-sampling [2] can handle piecewise-smooth constraints such as $\text{dist}(p(t, W), \mathcal{E}) \geq d_0$, they degrade the deterministic convergence guarantee to a probabilistic one in the sampling limit. In [15], [12], the authors note that two triangle meshes \mathcal{E} and \mathcal{E}' are d_0 -apart if and only if any line segment pair in $\mathcal{L}\mathcal{L}' \triangleq \mathcal{L} \times \mathcal{L}'$ and point-triangle pair in $\mathcal{PT}' \triangleq \mathcal{P} \times \mathcal{T}'$ or $\mathcal{P}'\mathcal{T} \triangleq \mathcal{P}' \times \mathcal{T}$ are d_0 -apart. As a result, the log-barrier of $\text{dist}(\bullet, \bullet)$ is equivalent to the sum of log-barrier of distance functions between primitive pairs, which can be made differentiable after an arbitrarily small perturbation.

We plan to use primal-only algorithms to solve UAV trajectory's local optimization problems. To this end, the three preconditions of the primal-only method must hold. The first condition of a feasible initial guess holds trivially by using an appropriate global searcher. With a fine-enough discretization granularity of signed distance field in [9], or that of sampled waypoints in [18], a feasible initial guess can be computed as long as a solution exists. For the second condition, although robust, exact continuous collision predicates exist for linear triangular meshes [1], [32], they have not been developed for polynomial curves. To bridge the gap, we propose to relax the exactness and resort to a conservative CCD scheme based on the subdivision of Bézier curve. For the third condition, i.e. the differentiability of the objective function, we propose to replace the trajectory with the union of convex hulls of the control polygons. These convex hulls are triangle meshes so the differentiable log-barrier between geometric primitives can be used. The accuracy of this approximation can be made arbitrarily high via adaptive subdivision.

B. P-IPM Framework

We use a column vector w to represent the set of control points of a Bézier curve. For the i -th piece Bézier curve of the trajectory, we denote its control points as $A_i W$, where A_i is a fixed transformation matrix of decision variables W into i -th curve's control points. A Bézier piece can be subdivided by a linear transformation of its control points into two pieces [7]: we denote control points of the first piece as $D_1 w$ and the control points of the second piece as $D_2 w$, where $D_{1,2}$ are the fixed subdivision stencils constructed using the De-Casteljau's algorithm. Our adaptive subdivision scheme (Algorithm 3) will recursively apply these stencils until a stopping criterion is met. And we keep a subdivision history \mathcal{H} to maintain the subdivision transformation matrices for each subdivided Bézier curve piece, where \mathcal{H} will be initialized with $\{A_i | i = 1, \dots, N\}$. Our algorithm handles two meshes \mathcal{E} and \mathcal{E}' . We use \mathcal{E} to denote the mesh of the environment and \mathcal{E}' denotes the mesh discretizing the UAV trajectory. After the subdivision, \mathcal{E}' will be updated so that it is the union of convex hulls of the new set of control points. For each subdivided Bézier curve piece with control points w , we denote $P(w)$ as the set of $M + 1$ vertices of the convex hull, $L(w)$ as the set of $(M + 1)M/2$ edges connecting all pair of vertices of $P(w)$, and $T(w)$ as

the set of $(M+1)M(M-1)/6$ triangles connecting any 3 vertices. As illustrated in Figure 1, given the environment \mathcal{E} (represented using either a point cloud or a triangle mesh) and \mathcal{E}' , we can define our log-barrier function as:

$$B(W, \mathcal{E}) \triangleq \sum_{A \in \mathcal{H}} \left[\sum_{l \in \mathcal{L}, l' \in L(w)} \text{clog}(\text{dist}(l, l') - d_0) + \sum_{p \in \mathcal{P}, t' \in T(w)} \text{clog}(\text{dist}(p, t') - d_0) + \sum_{t \in \mathcal{T}, p' \in P(w)} \text{clog}(\text{dist}(t, p') - d_0) \right] \Big|_{w=AW}, \quad (4)$$

where dist is the mollified, differentiable distance between a segment-segment or line-triangle pair, and clog is the clamped log-barrier function:

$$\text{clog}(x) = \begin{cases} -\frac{(x-x_0)^2}{x} \ln(\frac{x}{x_0}) & 0 < x \leq x_0 \\ 0 & \text{elsewhere,} \end{cases}$$

with x_0 being the activation range. Note that we use a slightly modified version of clog from the original paper [15]. Our modification preserves the second-order continuity but is essential for the correctness of inexact algorithm introduced in Section IV-D. Although there are many terms in Equation 4, all the terms with $\text{dist}(\bullet, \bullet)$ larger than $d_0 + x_0$ are zero. Therefore, we use a bounding volume hierarchy structure to efficiently prune all the zero terms.

Algorithm 1: NeedSub(w)

1: Return $\text{dist}(\text{Hull}(w), \mathcal{E}) < d_0 + x_0 \wedge \text{diameter}(\text{Hull}(w)) > \epsilon$

Algorithm 2: Sub(w, A, \mathcal{H})

1: $< \mathcal{P}', \mathcal{L}', \mathcal{T}' > \leftarrow < \emptyset, \emptyset, \emptyset >$
2: **if** NeedSub(w) **then**
3: $\mathcal{H} \leftarrow \mathcal{H} \setminus \{A\} \cup \{D_1A\} \cup \{D_2A\}$
4: $< \mathcal{P}', \mathcal{L}', \mathcal{T}' > \leftarrow < \mathcal{P}', \mathcal{L}', \mathcal{T}' > \cup \text{Sub}(D_1w, D_1A, \mathcal{H})$
5: $< \mathcal{P}', \mathcal{L}', \mathcal{T}' > \leftarrow < \mathcal{P}', \mathcal{L}', \mathcal{T}' > \cup \text{Sub}(D_2w, D_2A, \mathcal{H})$
6: **else**
7: $< \mathcal{P}', \mathcal{L}', \mathcal{T}' > \leftarrow < P(w), L(w), T(w) >$
8: Return $< \mathcal{P}', \mathcal{L}', \mathcal{T}' >$

Algorithm 3: TrajSub($p(t, W), \mathcal{H}$)

1: $< \mathcal{P}', \mathcal{L}', \mathcal{T}' > \leftarrow < \emptyset, \emptyset, \emptyset >$
2: **for** $A \in \mathcal{H}$ **do**
3: $< \mathcal{P}', \mathcal{L}', \mathcal{T}' > \leftarrow < \mathcal{P}', \mathcal{L}', \mathcal{T}' > \cup \text{Sub}(AW, A, \mathcal{H})$
4: Return $\mathcal{E}' \leftarrow < \mathcal{P}', \mathcal{L}', \mathcal{T}' >$

Algorithm 4: P-IPM

Input: W_0, T_0
1: $\mathcal{H} \leftarrow \{A_i | i = 1, \dots, N\}, W \leftarrow W_0, T \leftarrow T_0$
2: **while** True **do**
3: $\mathcal{E}' \leftarrow \text{TrajSub}(p(t, W), \mathcal{H})$
4: Calculate $\nabla_{W,T} O, \nabla_{W,T}^2 O$
5: Calculate $d \leftarrow -\text{SPD}(\nabla_{W,T}^2 O)^{-1} \nabla_{W,T} O$
6: $< W, T > \leftarrow \text{Search}(< W, T >, d)$
7: **if** $\|\nabla_{W,T} O\|_\infty \leq \epsilon_g$ **then**
8: Return $< W, T >$

We are now ready to summarize our main algorithm of P-IPM in Algorithm 4, which is a Newton-type method applied to the following unconstrained optimization problem:

$$\underset{W}{\text{argmin}} \quad O(W, T) + \lambda B(W, \mathcal{E}), \quad (5)$$

where λ is the relaxation coefficient controlling the exactness of constraint satisfaction. If the cost function is differentiable, the Hessian has bounded eigenvalues, then P-IPM converges to a first-order critical point of Equation 5. Finally, the function $\text{SPD}(\bullet)$ adjusts the Hessian matrix by clamping the negative eigenvalues to a small positive constant through a SVD to ensure positive-definiteness.

C. Convergence Guarantee

The convergence guarantee of Algorithm 4 relies on the correct implementation of two functions: Search (Algorithm 5) and NeedSub (Algorithm 1). The search function updates W to W' and ensures that the Wolfe's first condition:

$$O(W', T) + \lambda B(W', \mathcal{E}) \leq O(W, T) + \lambda B(W, \mathcal{E}) + c \alpha \nabla_W [O(W, T) + \lambda B(W, \mathcal{E})]^T d,$$

holds and $p(t, W)$ is homotopically equivalent to $p(t, W')$, where $c \in (0, 1)$ is some positive constant. Instead of using CCD which requires numerically stable polynomial root finding, we use a more conservative check between the convex hull of geometric primitives before and after the update from W to W' . This procedure is outlined in Algorithm 5.

The NeedSub function guarantees that the solution of Equation 5 asymptotically converge to that of Equation 3. For P-IPM, whenever there is an active log-barrier term in Equation 4, the convex hull of some Bézier curve with control points w is at most $d_0 + x_0$ away from \mathcal{E} . This implies we can bound the maximal distance between any curve's point and \mathcal{E} as $d_0 + x_0 + \Delta(w)$, where $\Delta(w)$ is the diameter of the convex hull of the control polygon or any upper bound between a curve's point and its convex hull. A simple strategy that guarantees ϵ -optimality is to always subdivide when 1) $\Delta(w) > \epsilon$ and 2) the distance between the convex hull of w and \mathcal{E} is smaller than $d_0 + x_0$, where ϵ is a user-provided optimality threshold. Note that condition 2) inherently induces an adaptive subdivision scheme where all the sub-trajectories that are sufficiently faraway from \mathcal{E} are not subdivided to save computation. Putting these components together, we show in our extended report [23] that, under mild assumptions, P-IPM will converge to the local optima of the original semi-infinite oracle (Equation 3) as the coefficients of log-barrier functions (λ) and the clamp range of clog functions (x_0) tends to zero.

Algorithm 5: Search($< W, T >, d$)

Input: $\gamma, c_1 \in (0, 1)$
1: **while** True **do**
2: $\alpha \leftarrow 1, < W', T' > \leftarrow < W, T > + \alpha d$
3: Safe \leftarrow True
4: **for** $A \in \mathcal{H}$ **do**
5: **if** $\text{dist}(\text{Hull}(AW \cup AW'), \mathcal{E}) < d_0$ **then**
6: Safe \leftarrow False
7: **if** Safe \wedge Wolfe's first condition **then**
8: Return $< W', T' >$
9: **else**
10: $\alpha \leftarrow \gamma \alpha, < W', T' > \leftarrow < W, T > + \alpha d$



Fig. 2: We optimize a curve involving a sharp turn in the middle using two different time-allocations. Left: Each curve is assigned the same amount of time. Right: The middle curve is assigned less time. For both cases, our optimizer generates similar trajectories (yellow dots are the end points of Bézier curves).

D. Inexact P-IPM

A practical problem with Algorithm 4 is that we have to add all the edges in $L(w)$ and all triangles in $T(w)$ of each convex hull of the control polygon to Equation 4. Even with a bounding volume hierarchy acceleration, summing up so many terms is still time-consuming. To overcome this issue, we propose an inexact, yet preserving the guarantees of local optimality, version of P-IPM. For every Bézier piece, we only keep all points and one edge. As compared with the exact counterpart, the inexact version significantly reduces the cost of log-barrier function evaluation by a factor of $O(M^3)$.

Note that most edges and all triangles are omitted in constructing Equation 4 but not in the Search function to ensure feasibility. This naive inexact P-IPM is not guaranteed to converge to the local minimum of Equation 5 because Algorithm 5 might not find a positive α satisfying the first Wolfe’s condition due to the conservative collision check. To ensure convergence to a local minimum, a simple strategy is to keep subdividing whenever the line-search fails. In our extended report [23], we show that, after several slight modifications to the line search Algorithm 5, the inexact P-IPM is guaranteed to converge to the KKT point after a finite number of subdivisions.

E. Time-Optimality

Time efficacy can be formulated as a terminal cost such as $R(p, T) = T$, combining the semi-infinite velocity and acceleration bounds. Unlike [30], [9] that warp each Bézier curve piece using a separate time parameter, we use a single, global time re-scaling. In other words, the UAV flies through each polynomial piece for T/N seconds. Although we use a single variable T as compared with N variables for each piece in [30], [9], we argue that time-optimality will not be sacrificed because our formulation allows a larger solution space for the relative length of the N Bézier curves. In prior work [9] for example, each Bézier curve is constrained to a separate convex subset of the freespace, which restricts the relative length of neighboring Bézier curves. By comparison, our method does not rely on these constraints and allow the relative length to change arbitrarily. As illustrated in Figure 2, we optimize two trajectories using different, fixed time allocations, and our method converges to almost identical solutions by changing the relative length of curves, which justifies the redundancy of curve-wise time variables.

We use the same principle as collision constraints to handle velocity and acceleration limits, by introducing a new set of log-barrier functions. For a subdivided curve piece with control points w , its first and second derivatives are two new Bézier curves, where their control points are

S_1w and S_2w , and $S_{1,2}$ are the fixed transformation matrix constructed by Bézier derivative. The convex hull of their control polygons and corresponding point, edge, triangle set are denoted as $\langle P(S_1w), L(S_1w), T(S_1w) \rangle$ for velocity and $\langle P(S_2w), L(S_2w), T(S_2w) \rangle$ for acceleration. We can now define our log-barrier function approximating the velocity and acceleration limits as:

$$B_T(W, T) = \sum_{i=1}^N \left[\sum_{p_1 \in P(S_1w)} \text{clog}(v_{max} - \|p_1\|_2 / (T/N)) + \sum_{p_2 \in P(S_2w)} \text{clog}(a_{max} - \|p_2\|_2 / (T^2/N^2)) \right]_{w=A_iW} . \quad (6)$$

P-IPM now minimizes $O(W, T) + \lambda B(W, \mathcal{E}) + \lambda B_T(W, T)$ to achieve joint optimality in terms of trajectory shape and time. Due to the convex hull property of Bézier curves, the finite value of $B_T(W, T)$ implies that Bézier curve with control points S_1w is bounded by Tv_{max}/N and Bézier curve with control points S_2w is bounded by T^2a_{max}/N^2 and subdivision can make the approximation arbitrarily exact. If desired, a separate subdivision rule can be used to control the exactness, e.g. always subdivide when $\Delta(S_1w) > \dot{\epsilon}$ or $\Delta(S_2w) > \ddot{\epsilon}$. Finally, note that we need a feasible initial guess for T and a sufficiently large T is always feasible.

V. EXPERIMENTS

Our implementation uses C++11 and all results are computed using a single thread on a workstation with a 3.5 GHz Intel Core i9 processor. Our algorithm has the following parameters: λ of the collision avoidance barrier term, x_0 of the activation range, d_0 of the clearance distance, v_{max} , a_{max} , ϵ of the subdivision threshold and ϵ_g of the stopping criterion. We use $\lambda = 10$, $x_0 = 0.1$, $d_0 = 0.1$, $v_{max} = 2.0m/s$, $a_{max} = 2.0m/s^2$, $\epsilon = 0.1$, $\epsilon_g = 10^{-3}$ for all experiments. To match the energy setting of the works we compare with, we use the composite Bézier where each piece is degree 8 and the continuity between adjacent pieces is C^2 , and set c in the objective function as jerk energy. We use improved GJK method [22] for convex hull collision detection.

A. Comparisons

We compare our approach with the state-of-the-art gradient based method [8] and corridor based method [9] on a set of scenes that are publicly downloadable from SketchFab [4]. We use the public implementations of both methods and tune our implementation to match the parameter settings of their methods, such as the orders of curves representing the trajectory, the energy term to measuring the smoothness of the trajectory (i.e. snap, jerk), v_{max} , and a_{max} . The implementation of [9] and [8] are run on dense point clouds sampled from the triangular meshes representing our scenes. Because of the complexity of the scenes listed in Table I, we manually tune waypoints to ensure that there will be a valid initial trajectory generated by the global search for all the methods, i.e. [9], [8], and ours.

As shown in Table I, our trajectories have the highest quality and far outperform the results of [9], [8]. Visual results are shown in Figure 3. Note that, since [8] doesn’t

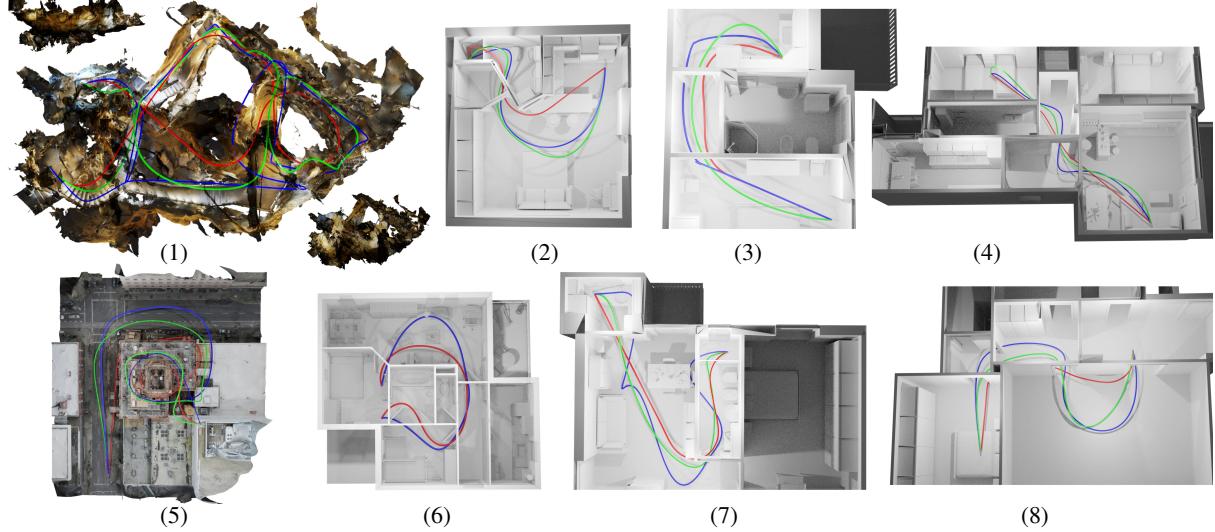


Fig. 3: For the scenes listed in Table I, we compare UAV trajectories computed using our method (red), [8] (blue), and [9] (green) from the same initial guess. Both our method and [9] can generate feasible trajectories and our method finds a smaller objective function (e.g. smoother trajectory and smaller travel time as shown in Table I), while [8] penetrates most of the environments.

Scene	Size	$L^{[9]}$	$T^{[9]}$	$C^{[9]}$	$L^{[8]}$	$T^{[8]}$	$C^{[8]}$	L^*	T^*	C^*
1	1.4M/2.7M	70.5	144.3	168.5	82.6	73.3	0.5	56.2	30.4	218.9
2	31K/58K	25.7	29.1	8.3	21.5	26.9	0.2	16.1	9.6	2.9
3	36K/68K	18.7	24.1	2.5	20.1	19.0	0.6	16.7	9.8	2.3
4	17K/34K	25.1	24.1	4.0	22.8	23.6	0.1	19.8	11.3	9.0
5	0.4M/0	49.8	65.0	6.4	50.0	50.0	0.3	38.2	21.3	42.1
6	0.2M/0.3M	-	-	-	14.8	18.0	0.2	13.8	9.2	14.9
7	81K/16K	23.2	38.5	3.8	30.9	26.3	0.2	21.6	12.4	6.2
8	36K/68K	24.5	23.5	3.3	28.1	26.1	0.1	18.5	10.7	2.2

TABLE I: We profile the performance of [9], [8], and our method (denoted as *) in terms of trajectory length L , arrival time T , and computational cost C . The size of each scene, in terms of its numbers of vertices and faces, is listed in the form of #V/#F. – denotes a failure in the provided program. Red color is used to highlight trajectories where the UAV collides with the scene.

have a collision-free guarantee even with a valid initial trajectory, their generated trajectories often penetrate the environments as highlighted in red in Table I. On the contrary, our method not only ensures a collision-free trajectory at every step during the optimization, but also guarantees that the trajectory is never closer to the obstacles than a user-specified safe distance. Although [9] requires only waypoints as input, we have to modify, in a trial-and-fail way, either the waypoints or the grid resolution used by their approach so that their method can find a collision-free initialization to ensure that their trajectory optimization proceeds.

Note that [8], [9] achieve real-time computations and good quality trajectories for simple environments, however they struggle as the environments gets more complex. [8] tends to generate invalid trajectories with collisions between the UAV and the environment. [9] instead has a quickly growing runtime, since they have to reduce the failure rate of their initial trajectory generation by increasing the grid resolution, resulting in a similar efficiency as ours.

B. Simulated UAV Flight in Challenging Environments

In the attached video (a screenshot is shown in Figure 4), we simulate the UAV flight trajectory in the challenging

environment of a museum with a lot of narrow passages. The simulator is setup using [21] on the hardware platform of a CrazyFlies nano-quadrotor. In the museum environment (represented using a dense triangle mesh with 1 million faces) we show that our optimization easily supports the generation of trajectories passing exactly through user-specified waypoints. As demonstrated in the video, our optimization can generate a high quality trajectory that has unnoticeable differences from the simulated path, allowing smooth UAV flying even during sharp turns.



Fig. 4: Left: a simulation screenshot of the UAV (red) flies along a simulated trajectory that has neglectable differences from our optimized trajectory (green) in a museum. Right: each plot shows the overlaid positions ($x/y/z[m]$) and the speeds ($xdot/ydot/zdot[m/s]$) of UAV for both the simulated and our trajectories.

VI. CONCLUSION & LIMITATIONS

We propose a new approach of local UAV-trajectory optimization with guaranteed feasibility and asymptotic convergence to the semi-infinite problem. The key to our success is the primal-only optimizer equipped with an adaptive subdivision scheme for line search and constraint refinement. The line search scheme preserves the homotopy class and the constraint refinement guarantees sufficient closedness to the semi-infinite oracle (Equation 3). Using the same framework, we show that our method can be extended to achieve time-optimality. Besides, our method requires a feasible initial guess, which can be easily computed using global search techniques.

REFERENCES

- [1] T. Brochu, E. Edwards, and R. Bridson, "Efficient geometrically exact continuous collision detection," *ACM Trans. Graph.*, vol. 31, no. 4, July 2012. [Online]. Available: <https://doi.org/10.1145/2185520.2185592>
- [2] F. E. Curtis and M. L. Overton, "A sequential quadratic programming algorithm for nonconvex, nonsmooth constrained optimization," *SIAM Journal on Optimization*, vol. 22, no. 2, pp. 474–500, 2012.
- [3] R. Deits and R. Tedrake, "Efficient mixed-integer planning for uavs in cluttered environments," in *2015 IEEE international conference on robotics and automation (ICRA)*. IEEE, 2015, pp. 42–49.
- [4] A. Denoyel, C. Pinson, and P.-A. Passet, *SketchFab*, 2020 (accessed October 25, 2020). [Online]. Available: <https://sketchfab.com/>
- [5] C. Di Franco and G. Buttazzo, "Energy-aware coverage path planning of uavs," in *2015 IEEE international conference on autonomous robot systems and competitions*. IEEE, 2015, pp. 111–117.
- [6] W. Ding, W. Gao, K. Wang, and S. Shen, "Trajectory replanning for quadrotors using kinodynamic search and elastic optimization," in *2018 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, 2018, pp. 7595–7602.
- [7] G. E. Farin and D. Hansford, *The Essentials of CAGD*, 1st ed. USA: A. K. Peters, Ltd., 2000.
- [8] F. Gao, Y. Lin, and S. Shen, "Gradient-based online safe trajectory generation for quadrotor flight in complex environments," in *2017 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2017, pp. 3681–3688.
- [9] F. Gao, W. Wu, Y. Lin, and S. Shen, "Online safe trajectory generation for quadrotors using fast marching method and bernstein basis polynomial," in *2018 IEEE International Conference on Robotics and Automation (ICRA)*, 2018, pp. 344–351.
- [10] D. Grundel and D. Jeffcoat, "Formulation and solution of the target visitation problem," in *AIAA 1st Intelligent Systems Technical Conference*, 2004, p. 6212.
- [11] J. A. Guerrero and Y. Bestaoui, "Uav path planning for structure inspection in windy environments," *Journal of Intelligent & Robotic Systems*, vol. 69, no. 1-4, pp. 297–311, 2013.
- [12] K. Hauser, "Semi-infinite programming for trajectory optimization with non-convex obstacles," in *International Workshop on the Algorithmic Foundations of Robotics*. Springer, 2018, pp. 565–580.
- [13] M. Kalakrishnan, S. Chitta, E. Theodorou, P. Pastor, and S. Schaal, "Stomp: Stochastic trajectory optimization for motion planning," in *2011 IEEE international conference on robotics and automation*. IEEE, 2011, pp. 4569–4574.
- [14] S. Kim, K. Sreenath, S. Bhattacharya, and V. Kumar, "Optimal trajectory generation under homology class constraints," in *2012 IEEE 51st IEEE Conference on Decision and Control (CDC)*, 2012, pp. 3157–3164.
- [15] M. Li, Z. Ferguson, T. Schneider, T. Langlois, D. Zorin, D. Panozzo, C. Jiang, and D. M. Kaufman, "Incremental potential contact: Intersection-and inversion-free, large-deformation dynamics," *ACM Trans. Graph.*, vol. 39, no. 4, July 2020. [Online]. Available: <https://doi.org/10.1145/3386569.3392425>
- [16] Y. Li, Z. Littlefield, and K. E. Bekris, "Asymptotically optimal sampling-based kinodynamic planning," *The International Journal of Robotics Research*, vol. 35, no. 5, pp. 528–564, 2016.
- [17] M. Likhachev, D. Ferguson, G. Gordon, A. Stentz, and S. Thrun, "Anytime dynamic a*: An anytime, replanning algorithm," in *Proceedings of the Fifteenth International Conference on International Conference on Automated Planning and Scheduling*, ser. ICAPS'05. AAAI Press, 2005, p. 262–271.
- [18] S. Liu, N. Atanasov, K. Mohta, and V. Kumar, "Search-based motion planning for quadrotors using linear quadratic minimum time control," in *2017 IEEE/RSJ international conference on intelligent robots and systems (IROS)*. IEEE, 2017, pp. 2872–2879.
- [19] O. Mechali, L. Xu, M. Wei, I. Benkhadra, F. Guo, and A. Senouci, "A rectified rrt* with efficient obstacles avoidance method for uav in 3d environment," in *2019 IEEE 9th Annual International Conference on CYBER Technology in Automation, Control, and Intelligent Systems (CYBER)*. IEEE, 2019, pp. 480–485.
- [20] D. Mellinger and V. Kumar, "Minimum snap trajectory generation and control for quadrotors," in *2011 IEEE International Conference on Robotics and Automation*, 2011, pp. 2520–2525.
- [21] N. Michael, D. Mellinger, Q. Lindsey, and V. Kumar, "The grasp multiple micro-uav testbed," *IEEE Robotics Automation Magazine*, vol. 17, no. 3, pp. 56–65, 2010.
- [22] M. Montanari, N. Petrinic, and E. Barbieri, "Improving the gjk algorithm for faster and more reliable distance queries between convex objects," *ACM Transactions on Graphics (TOG)*, vol. 36, no. 3, pp. 1–17, 2017.
- [23] R. Ni, T. Schneider, D. Panozzo, Z. Pan, and X. Gao, "Robust & asymptotically locally optimal uav-trajectory generation based on spline subdivision," *arXiv preprint arXiv:2010.09904*, 2020.
- [24] C. Park, J. Pan, and D. Manocha, "Itomp: Incremental trajectory optimization for real-time replanning in dynamic environments," in *Twenty-Second International Conference on Automated Planning and Scheduling*, 2012.
- [25] Q.-C. Pham, "A general, fast, and robust implementation of the time-optimal path parameterization algorithm," *IEEE Transactions on Robotics*, vol. 30, no. 6, pp. 1533–1540, 2014.
- [26] H. Prautzsch, W. Boehm, and M. Paluszny, *Bezier and B-Spline Techniques*. Berlin, Heidelberg: Springer-Verlag, 2002.
- [27] C. Richter, A. Bry, and N. Roy, "Polynomial trajectory planning for aggressive quadrotor flight in dense indoor environments," in *Robotics Research*. Springer, 2016, pp. 649–666.
- [28] J. Schulman, Y. Duan, J. Ho, A. Lee, I. Awwal, H. Bradlow, J. Pan, S. Patil, K. Goldberg, and P. Abbeel, "Motion planning with sequential convex optimization and convex collision checking," *The International Journal of Robotics Research*, vol. 33, no. 9, pp. 1251–1270, 2014.
- [29] N. Slegers, J. Kyle, and M. Costello, "Nonlinear model predictive control technique for unmanned air vehicles," *Journal of guidance, control, and dynamics*, vol. 29, no. 5, pp. 1179–1188, 2006.
- [30] W. Sun, G. Tang, and K. Hauser, "Fast uav trajectory optimization using bilevel optimization with analytical gradients," in *2020 American Control Conference (ACC)*. IEEE, 2020, pp. 82–87.
- [31] G. Tang, W. Sun, and K. Hauser, "Enhancing bilevel optimization for uav time-optimal trajectory using a duality gap approach," in *2020 IEEE International Conference on Robotics and Automation (ICRA)*, 2020, pp. 2515–2521.
- [32] M. Tang, R. Tong, Z. Wang, and D. Manocha, "Fast and exact continuous collision detection with bernstein sign classification," *ACM Transactions on Graphics (TOG)*, vol. 33, no. 6, pp. 1–8, 2014.
- [33] O. Von Stryk, "Numerical solution of optimal control problems by direct collocation," in *Optimal control*. Springer, 1993, pp. 129–143.
- [34] Z. Wang, X. Zhou, C. Xu, J. Chu, and F. Gao, "Alternating minimization based trajectory generation for quadrotor aggressive flight," *IEEE Robotics and Automation Letters*, vol. 5, no. 3, pp. 4836–4843, 2020.
- [35] D. J. Webb and J. van den Berg, "Kinodynamic rrt*: Optimal motion planning for systems with linear differential constraints," *arXiv*, pp. arXiv-1205, 2012.
- [36] D. J. Webb and J. Van Den Berg, "Kinodynamic rrt*: Asymptotically optimal motion planning for robots with linear dynamics," in *2013 IEEE International Conference on Robotics and Automation*. IEEE, 2013, pp. 5054–5061.
- [37] G. Williams, P. Drews, B. Goldfain, J. M. Rehg, and E. A. Theodorou, "Aggressive driving with model predictive path integral control," in *2016 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, 2016, pp. 1433–1440.
- [38] B. Zhou, F. Gao, L. Wang, C. Liu, and S. Shen, "Robust and efficient quadrotor trajectory generation for fast autonomous flight," *IEEE Robotics and Automation Letters*, vol. 4, no. 4, pp. 3529–3536, 2019.
- [39] M. Zucker, N. Ratliff, A. D. Dragan, M. Pivtoraiko, M. Klingensmith, C. M. Dellin, J. A. Bagnell, and S. S. Srinivasa, "Chomp: Covariant hamiltonian optimization for motion planning," *The International Journal of Robotics Research*, vol. 32, no. 9-10, pp. 1164–1193, 2013.